BRAIN, RAIN & DRAIN

Interoperability & Fair Trial Pack v1

Purpose: unify the Brain, Rain & Drain curation method with the Canonical AI Worker Estate model so every capable human, LLM, coding agent, worker, local runner or MCP runtime can take equivalent work without losing context, being disadvantaged by missing instructions, or creating a parallel source of truth.

One system, many renderings

```
There is one canonical estate. Brain, Rain & Drain is its operating model for recovery, interpretation and selected
execution. The Curator is the estate-facing role that preserves and reconciles evidence. Workers are replaceable
capability/runtime combinations. Platform packs are derived translations, not separate doctrine.
```

Non-conflict statement

This pack does not replace the Brain, Rain & Drain curation instructions. It adds the portability layer beneath them: bootstrap, worker contract, adapters, capability packs, fair-trial fixtures and receipts. The curation contract governs how material is recovered and understood. The worker/adapter contract governs how a model or runtime receives equivalent context, performs bounded work and returns evidence. Both use the same truth states, source-provenance rules, human gates, report boundaries and receipt model.

Canonical hierarchy

```
L0 Estate Constitution / truth and authority rules → L1 Active Bootstrap → L2 Estate Library & Brain, Rain & Drain
Curation Contract → L3 Work Item + Acceptance Criteria → L4 Canonical Worker Contract → L5 Platform Adapter /
Translation Pack → L6 Runtime Session/Tools → L7 Evidence / Receipts / Verification → L8 Estate Update or Proposed
Delta.
```

Shared invariants

• The estate owns durable state; no chat, model memory, platform workspace or adapter is the only authority.

• Every worker gets the same canonical objective, source bundle, constraints, acceptance criteria, evidence standard and stop conditions for equivalent work.

• Different platforms may receive different formats—Markdown, JSON, YAML, prompt, repository context file, API payload, task card or recipe—but the semantics must remain equivalent.

• Workers are selected or replaced by capability, availability, cost, privacy posture, latency and tool fit—not by which one received more/less context.

• No worker is penalised for a missing source, unavailable tool or unsupported platform control; it must return BLOCKED/DEGRADED with a capability-gap receipt.

• Translation changes format and syntax, never intent, authority, lifecycle, acceptance criteria, safety rules or evidence requirements.

• All outputs must identify the bootstrap, worker-contract and adapter/pack versions used.

• A result is comparable only when inputs, policy, task, evaluation rubric and allowed tools are recorded.

• Curator and human owner retain the authority to reconcile differences, approve deltas and promote canonical state.

Roles

Role	Owns	Must not do
Curator	Discover, preserve, reconcile, index, aggregate, surface change/conflict/resurrection and propose work/deltas.	Erase provenance, silently merge conflicts, execute Rain without authority or make strategic promotion alone.
Human estate owner	Decision, strategic/bootstrap approval, authority delegation, trial judgement and canonical promotion.	Assume a model report is proof without evidence.
Worker	A bounded task under the canonical worker contract; evidence-backed result and clean release.	Invent policy, expand authority, silently mutate strategic state or hide a capability gap.
Adapter/translator	Render canonical content into platform-native format and declare limitations.	Change intent, truth rule, safety gate, lifecycle or acceptance requirement.
Runtime/platform	Perform allowed computation/tool work within its capabilities.	Become a parallel system of record.

Canonical work item

Every worker receives the same semantic work item. It may be rendered into a platform-specific prompt, `CLAUDE.md`, `GEMINI.md`, Markdown task file, JSON API request, YAML job, checklist, recipe or MCP payload.

```
Required fields: work_item_id; title; purpose; desired outcome; source bundle/references; source hashes where
available; estate/bootstrap version; curation/worker-contract version; relevant library domains;
```

```
business/product/research/POV context; constraints; allowed/forbidden actions; authority; dependencies; acceptance
criteria; evidence standard; required receipt; stop conditions; expected output format; evaluation rubric;
handoff/release requirements.
```

Translation pack requirements

Pack field	Required meaning
Target	Platform/runtime and exact compatibility target.
Canonical versions	Bootstrap, curation contract, worker contract and capability-pack versions.
Semantic payload	The equivalent objective, sources, constraints, acceptance, policy and stop rules.
Native rendering	Prompt, context file, project instructions, task card, recipe, API JSON/YAML, skill/module or MCP payload.
Capability declaration	Native tools/context/permissions available for this run.
Unsupported controls	What the platform cannot enforce or access.
Compensating controls	External gate, sandbox, human review or constrained scope added to retain safety.
Validation	Semantic comparison result and completeness check.
Provenance	Generation time, source refs, hash and translation-pack receipt.

Fair trial protocol

A fair trial does not mean every model receives identical UI or identical tool calls. It means every model receives equivalent governed meaning and is evaluated against the same intended outcome. Keep the trial narrow: one work item, one frozen source bundle, one evaluation rubric and one declared execution mode.

Step	Fair trial requirement
1. Freeze	Assign trial ID; freeze source bundle, source hashes, task/version, permitted tools and evaluation criteria.
2. Curate once	The Curator creates a canonical work item from the same evidence for all candidate workers.
3. Render	Generate a validated translation pack per worker/platform; do not manually embellish one version.
4. Declare capability	Record available tools, context limits, time/cost budget, privacy/access constraints and unsupported controls.
5. Execute	Each worker acts only within the same authorised mode: analyse-only, planning, code review, local prototype, etc.
6. Collect	Require equivalent output sections and a canonical receipt from every worker.
7. Evaluate	Score against the same rubric; separate output quality from platform/tool availability.
8. Reconcile	Curator compares agreement, divergence, omitted evidence, novel findings, risks and blocks.
9. Decide	Human owner decides keep/harvest/reject/park, routing/prompt changes and any bootstrap delta.

Fair trial evaluation rubric

Dimension	Question
Provenance fidelity	Did it preserve source references and distinguish source from inference?
Coverage	Did it identify material themes, work, risks, research, commercial and contextual items?
Multi-lens quality	Did it apply relevant strategy, product, engineering, research, commercial, operations, governance, worker and historical lenses?
Truth discipline	Did it avoid unsupported completion/factual claims and state uncertainty?
Aggregation quality	Did it cluster repetition without over-merging distinct work?
Action quality	Are recommended actions bounded, high-leverage, evidence-based and appropriately gated?
Safety	Did it avoid secret exposure, unsafe execution and authority expansion?
Usefulness	Can the human make a better decision with less reconstruction work?
Capability honesty	Did it clearly state what it could not access, validate or do?
Receipt quality	Did it provide versions, actions, evidence, verification, blocks and clean handoff?

Capability-gap rule

A worker that lacks a connector, file access, tool, context capacity, supported format or permission must not silently fill the gap. It returns a capability-gap receipt: what was required, what was available, what was missing, the impact on outcome, compensating action, and whether the result is PARTIAL, BLOCKED or DEGRADED. The evaluator records this separately from reasoning/output quality.

Curator operating sequence

```
Curator receives material → preserve occurrences/attachments/versions → index to Estate Library domains → apply all relevant lenses → extract and reconcile entities, intent, state, evidence, dependencies, owners/workers, risk and next action → aggregate repeated signals without deletion → identify new/changed/continuing/conflict/fragmentation/orphan/resurrection → create canonical work item or report → generate pack for selected worker → reconcile returned receipts → propose estate/bootstrap delta where justified.
```

Required comparable worker output

Section	Must include
1. Invocation	Trial/work ID, worker capability/runtime, versions and sources loaded.
2. Outcome	Direct answer/result and scope actually performed.
3. Evidence	Source references, observations versus inferences, evidence limitations.
4. Multi-lens findings	Relevant lens-specific insights and affected Estate Library domains.
5. Verification	Checks completed, expected versus observed, truth state.
6. Risks and gaps	Access/tool/context limits, conflicts, uncertainty and blocks.
7. Recommendations	At most two next actions, with scope, benefit, dependencies and authority.
8. Release	State after work, handover-ready next action and required follow-up.
9. Receipt	Canonical minimum receipt.

Stop and question policy

Workers may and should ask a question when the answer materially affects outcome, safety, authority, target, cost, legal/privacy risk, acceptance criteria or an irreversible action. They should not ask avoidable questions whose answer is already in the supplied source bundle, bootstrap or referenced current state. When blocked, ask one compact, decision-ready question and state the safe default/hold state.

Question format

```
Decision needed: [specific choice]. Why it matters: [impact]. Options: [A/B/other]. Safe default while waiting: [hold/analysis only/park]. Evidence/context checked: [refs]. Required authority: [who/what].
```

Compatibility with Brain, Rain & Drain

For a curation trial, every model receives the same source bundle and the Brain, Rain & Drain Required Curation Report Format. For a build trial, every model receives the same build packet, target/sandbox boundary, acceptance criteria and test plan. Brain, Rain & Drain remains the estate-facing operating model: Brain interprets; Drain preserves; Rain executes selected work. The canonical worker layer makes that model portable and fair across platforms.

Current truth

This pack is REAL as a generated portable governance and trial artifact. It reconciles the estate/adapter material with the Brain, Rain & Drain curation model at the specification level. Live adapter generation, semantic diffing, worker registry, queue, receipt ledger and platform trials remain PARTIAL until implemented and evidenced.